

## Q1. Program: Create soft link, hard link, and FIFO in Linux

### 1. Objective

The aim is to create different types of files in Linux using **system calls**:

- **Soft link (symbolic link)**: Like a shortcut → points to the original file's path.
- **Hard link**: Another name for the same file → shares the same inode.
- **FIFO (named pipe)**: A special file used for **inter-process communication (IPC)**.

We will also see how to create them using **shell commands** vs **system calls**.

### 2. Header Libraries (and why)

```
#include <stdio.h>    // For printf(), fprintf(), fopen(), fclose()  
  
#include <unistd.h>   // For symlink(), link()  
  
#include <sys/stat.h> // For mkfifo(), file permission modes
```

### 3. Equivalent Shell Commands

(i) Soft link

```
ln -s original.txt soft_link.txt
```

(ii) Hard link

```
ln original.txt hard_link.txt
```

(iii) FIFO

```
mkfifo myfifo
```

### 4. Full Code (with comments)

```
#include <stdio.h>  
  
#include <unistd.h>   // symlink(), link()  
  
#include <sys/stat.h> // mkfifo()  
  
int main() {  
  
    // 1. Create a normal file  
  
    FILE *fp = fopen("original.txt", "w");  
  
    fprintf(fp, "This is the original file.\n");  
  
    fclose(fp);
```

```

// 2. Create a symbolic (soft) link

symlink("original.txt", "soft_link.txt");

printf("Soft link created: soft_link.txt -> original.txt\n");

// 3. Create a hard link

link("original.txt", "hard_link.txt");

printf("Hard link created: hard_link.txt -> original.txt\n");

// 4. Create a FIFO (named pipe)

mkfifo("myfifo", 0666);

printf("FIFO created: myfifo\n");

return 0;
}

```

## 5. Step-by-step Explanation

### 1. Normal file creation

- `fopen("original.txt", "w")`: Creates `original.txt`.
- `fprintf()` writes some content.
- `fclose()` closes file.

### 2. Soft link

- `symlink("original.txt", "soft_link.txt")`: Creates a symbolic link.
- `soft_link.txt` stores the **path** of `original.txt`.
- If `original.txt` is deleted, the soft link becomes **dangling** (broken).

### 3. Hard link

- `link("original.txt", "hard_link.txt")`: Creates a hard link.
- Both files share the **same inode number**.
- Even if `original.txt` is deleted, `hard_link.txt` still contains the data.

#### 4. FIFO

- `mkfifo("myfifo", 0666)`: Creates a named pipe.
  - It's a special file for **process-to-process communication**.
  - 0666: Read & write permission for everyone.
- 

#### 6. Sample Output

When program runs:

Soft link created: `soft_link.txt -> original.txt`

Hard link created: `hard_link.txt -> original.txt`

FIFO created: `myfifo`

Then check with:

`ls -l`

Example:

`-rw-r--r-- 1 user user 28 Sep 13 21:30 original.txt`

`lrwxrwxrwx 1 user user 12 Sep 13 21:30 soft_link.txt -> original.txt`

`-rw-r--r-- 2 user user 28 Sep 13 21:30 hard_link.txt`

`prw-rw-rw- 1 user user 0 Sep 13 21:30 myfifo`

---

#### 7. Viva Q&A (what examiner may ask)

- **Q: Difference between soft and hard link?**

**A:** *Soft link*: Points to filename (can cross partitions). Broken if original file is deleted.

*Hard link*: Points to same inode (cannot cross partitions). Survives even if original is deleted.

- **Q: What is FIFO used for?**

**A:** FIFO (named pipe) is used for **IPC** so processes can read/write like a normal file but in a synchronized manner.

- **Q: Which system calls are used?**

**A:** `symlink()`, `link()`, `mkfifo()`.

- **Q: How is FIFO different from an unnamed pipe?**

**A:** Unnamed pipe exists only between related processes and disappears when closed. FIFO is a special file visible in filesystem and can be used by unrelated processes.

---

## 8. Common Mistakes

- Forgetting to include <sys/stat.h> → mkfifo will not compile.
- Not checking return values of system calls → can fail if file already exists.
- Running multiple times → second run may fail because soft\_link.txt, hard\_link.txt, myfifo already exist. Use unlink() before recreating.

## Q2: Infinite Loop & Process Info in /proc

### 1. Objective

- Write a process that runs continuously in the background.
- Then go to /proc/<PID>/ to explore all information about that process (memory, state, threads, permissions, etc.).
- This helps us understand how **Linux manages processes** and exposes details through the /proc virtual filesystem.

---

### 2. Header Libraries

```
#include <stdio.h> // printf()  
#include <unistd.h> // getpid(), sleep()
```

---

### 3. Code (with comments)

```
#include <stdio.h>  
  
#include <unistd.h>  
  
int main() {  
    pid_t pid = getpid(); // get process ID of this program  
    printf("Process started. PID: %d\n", pid);  
  
    // infinite loop: keep process alive  
    while (1) {  
        printf("Process %d is running\n", pid);  
        sleep(2); // sleep 2 seconds to avoid CPU hogging  
    }  
    return 0;  
}
```

## 4. Step-by-Step Explanation

1. **getpid()** → returns current process ID (PID). This PID is used to look up details in `/proc/<pid>/`.
  2. **Infinite loop (while(1))** → keeps the process alive. Without this, process would exit immediately.
  3. **printf inside loop** → shows process is running.
  4. **sleep(2)** → makes the process wait for 2 seconds between prints (saves CPU).
- 

## 5. Running the Program

```
gcc 2.c -o myprog
```

```
./myprog &
```

- & runs the program in the background.
  - You will see the **PID** printed, e.g., PID: 422.
- 

## 6. Checking Process in /proc

- `/proc` is a **virtual filesystem** that contains a directory for each running process, named by its PID.
- Example:
- `cat /proc/422/status`

will show all info about PID **422**.

---

## 7. Understanding `/proc/<pid>/status` Output

Some important fields you should mention in viva:

- **Name:** Executable name (`a.out` in this case).
- **Pid:** Process ID.
- **PPid:** Parent process ID.
- **State:** Current state (R running, S sleeping, Z zombie, etc.).
- **Uid / Gid:** User and group IDs that own the process.
- **VmSize, VmRSS, VmData:** Memory usage (Virtual Memory, Resident Set Size, Data segment).
- **Threads:** Number of threads used.
- **FDSIZE:** Number of file descriptors allocated.

- **voluntary\_ctxt\_switches / nonvoluntary\_ctxt\_switches:** How many times the process switched context voluntarily or due to scheduler.
- 

## 8. Viva Q&A

- **Q: What is /proc?**  
**A:** /proc is a virtual filesystem in Linux that provides kernel and process information in file form. It doesn't exist on disk; it's generated by the kernel dynamically.
- **Q: What is difference between /proc/<pid>/status and /proc/<pid>/stat?**  
**A:** status: Human-readable summary.  
  
stat: Raw numbers used by system tools like ps.
- **Q: What happens if you kill the process?**  
**A:** Its /proc/<pid>/ directory disappears immediately, since /proc reflects live processes.
- **Q: What is process state S?**  
**A:** Sleeping — process is waiting (e.g., due to sleep() system call).
- **Q: How to run process in background?**  
**A:** Append & after command (e.g., ./a.out &).

## Q3. Program: Create a File and Print File Descriptor using creat()

### 1. Objective

Write a program that:

- Creates a new file using the **creat() system call**.
  - Prints the **file descriptor value** returned.
  - Demonstrates file permissions (0644) and concept of file descriptors.
- 

### 2. Required Header Libraries

- **<stdio.h>** → for printf(), perror().
- **<fcntl.h>** → contains declaration of creat() and file permission flags.
- **<unistd.h>** → for close() system call.

### 3. Code

```
#include <stdio.h>
#include <fcntl.h>
#include <unistd.h>

int main() {
    // Create a file named "myfile.txt" with permission 0644
```

```

int fd = creat("myfile.txt", 0644);

// Print file descriptor value

printf("File created successfully, descriptor value: %d\n", fd);

// Close the file

close(fd);

return 0;

}

```

---

#### 4. Step-by-Step Explanation

##### 1. `creat("myfile.txt", 0644)`

- Creates a file myfile.txt.
- If it already exists, it truncates (empties) the file.
- 0644 = permission mode (rw-r--r--).

##### 2. Return value

- On success → returns a **file descriptor** (non-negative int).
- On failure → returns -1.

##### 3. File Descriptor (FD)

- An integer index into the kernel's open file table.
- Standard FDs:
  - 0 → stdin
  - 1 → stdout
  - 2 → stderr
- First new file → usually gets 3.

##### 4. `close(fd)`

- Releases the file descriptor.

---

#### 5. Compilation & Execution

`gcc 3.c -o filecreate`  
`./filecreate`

## Sample Output

File created successfully, descriptor value: 3

Check file:

ls -l myfile.txt

Output:

-rw-r--r-- 1 user user 0 Sep 13 22:10 myfile.txt

---

## 6. Viva Questions & Answers

- **Q: What is creat() used for?**

**A:** To create a file (or truncate if it exists). Equivalent to open() with flags O\_CREAT | O\_WRONLY | O\_TRUNC.

- **Q: What is a file descriptor?**

**A:** A non-negative integer returned by the kernel that uniquely identifies an open file within a process.

- **Q: Why 0644?**

**A:** Permission bits in octal. It means → Owner: read+write, Group: read, Others: read.

- **Q: What are file descriptors 0, 1, 2?**

**A:** stdin, stdout, stderr.

- **Q: How is creat() different from open()?**

**A:** creat() is a simpler wrapper; internally it just calls open() with fixed flags.

## Q4. Write a program to open an existing file with read write mode. Try O\_EXCL flag also.

### Program Goal

- Demonstrate how to **open/create files in read–write mode.**
  - Show the effect of using the **O\_EXCL flag** with open().
- 

### Header Files

```
#include <stdio.h> // for printf()  
  
#include <fcntl.h> // for open() and file control flags (O_RDWR, O_CREAT, O_EXCL)  
  
#include <unistd.h> // for close()
```

---

### System Call: open()

```
int open(const char *pathname, int flags, mode_t mode);
```

- **pathname** → name of the file (e.g., "txtFile4.txt").
- **flags** → specify how the file should be opened.
- **mode** → used only if the file is created (defines permissions like 0644).

Returns:

- File Descriptor (FD)  $\geq 0$  on success.
  - -1 on failure.
- 

## Flags Used

1. **O\_RDWR** → Open file for **both reading and writing**.
  2. **O\_CREAT** → Create the file if it does not already exist.
    - Needs mode (here 0644 → rw-r--r--).
  3. **O\_EXCL** → Used with O\_CREAT.
    - If the file already exists → **open fails**.
    - If the file does not exist → file is created.
- 

## Code Explanation

### Part 1:

```
fd = open("txtFile4.txt", O_RDWR | O_CREAT, 0644);
printf("file created/opened. FD: %d\n", fd);
close(fd);
```

- Opens txtFile4.txt in read–write mode.
  - If file doesn't exist → creates it with permissions 0644 (owner can read/write, others can read).
  - Prints the FD (e.g., 3).
- 

### Part 2:

```
fd = open("txtFile.txt", O_CREAT | O_EXCL | O_RDWR, 0644);
printf("O_EXCL succeeded. FD = %d\n", fd);
close(fd);
```

- Tries to create txtFile.txt with O\_EXCL.
- If file **already exists** → this call **fails**.

- If file does **not exist** → creates successfully and prints FD.
- 

## Output Example

file created/opened. FD: 3

O\_EXCL succeeded. FD = 3

- First line: txtFile4.txt created/opened.
- Second line: txtFile.txt created since it didn't exist before.

If you run again without deleting txtFile.txt, the second open() would fail (because of O\_EXCL).

---

## Viva Questions

1. **Q:** What is the purpose of O\_EXCL?

**A:** It ensures the call fails if the file already exists, preventing accidental overwrite.

2. **Q:** What is the difference between creat() and open() with O\_CREAT?

**A:** creat() is just a simplified version of open() that always uses O\_CREAT | O\_WRONLY | O\_TRUNC.

3. **Q:** What does 0644 mean?

**A:** File permissions → Owner: read/write, Group: read, Others: read.

## Q5. Write a program to create five new files with infinite loop. Execute the program in the background and check the file descriptor table at /proc/pid/fd.

### 1. Creating multiple files

```
int fd[5];
for (int i = 0; i < 5; i++) {
    char filename[20];
    sprintf(filename, "file%d.txt", i + 1);
    fd[i] = open(filename, O_CREAT | O_RDWR, 0644);
    printf("created %s with fd= %d\n", filename, fd[i]);
}
```

- An array fd[5] stores the **file descriptors** of the 5 files.
- Loop runs from 0 to 4.
- Uses sprintf() to generate file names:
  - file1.txt, file2.txt, ..., file5.txt.

- Each file is opened with:
    - O\_CREAT → create file if not present.
    - O\_RDWR → open in read-write mode.
    - 0644 → file permission rw-r--r--.
  - The returned file descriptor (an integer) is printed.
- 

## 2. Printing Process ID

```
printf("process ID: %d\n", getpid());
```

- getpid() returns the **process ID** of the running program.
  - Important because we'll need it to check /proc/<pid>/fd.
- 

## 3. Infinite Loop

```
while (1){  
    // keeps the program alive  
}
```

- Keeps the process running in background.
  - Ensures files remain open so we can inspect /proc/<pid>/fd.
- 

## How to Run

1. Compile:
2. gcc 5.c -o q5
3. Run in background:
4. ./q5 &

Example output:

created file1.txt with fd= 3

created file2.txt with fd= 4

created file3.txt with fd= 5

created file4.txt with fd= 6

created file5.txt with fd= 7

process ID: 5123

## Checking File Descriptor Table

Go to process directory:

```
cd /proc/5123/fd
```

```
ls -l
```

Output (example):

```
0 -> /dev/pts/0 (stdin)
```

```
1 -> /dev/pts/0 (stdout)
```

```
2 -> /dev/pts/0 (stderr)
```

```
3 -> /home/jils/file1.txt
```

```
4 -> /home/jils/file2.txt
```

```
5 -> /home/jils/file3.txt
```

```
6 -> /home/jils/file4.txt
```

```
7 -> /home/jils/file5.txt
```

- FDs **0,1,2** are standard input/output/error.
- FDs **3–7** correspond to file1.txt ... file5.txt.

---

## Viva Questions

1. **Q:** What are file descriptors 0, 1, 2?  
**A:** 0 = stdin, 1 = stdout, 2 = stderr.
2. **Q:** Why use an infinite loop here?  
**A:** To keep process alive, so we can check /proc/<pid>/fd.
3. **Q:** What happens if we close a file descriptor?  
**A:** Its entry disappears from /proc/<pid>/fd.
4. **Q:** How does Linux represent open files in /proc/<pid>/fd?  
**A:** As **symbolic links** pointing to the actual file.

## Q6. Write a program to take input from STDIN and display on STDOUT. Use only read/write system calls

### Program Goal

- Take input from **standard input (stdin)** and display it on **standard output (stdout)**.
- Use only **low-level system calls**: read() and write().
- No use of printf() or scanf() (high-level C I/O).

## Header Files

```
#include <unistd.h> // for read() and write() system calls
```

- `read()` → reads data from a file descriptor.
  - `write()` → writes data to a file descriptor.
- 

## Code (Consolas style)

```
#include <unistd.h>

int main() {
    char c;
    // Infinite loop to read one character at a time from stdin
    while (read(0, &c, 1) > 0) {
        // Write the character to stdout
        write(1, &c, 1);
    }
    return 0;
}
```

---

## Step-by-Step Explanation

### 1. `read(0, &c, 1)`

- 0 → file descriptor for **stdin**.
- `&c` → address of variable `c` where the input character will be stored.
- 1 → read **1 byte at a time**.
- Returns the number of bytes read (>0) or 0 at EOF.

### 2. `write(1, &c, 1)`

- 1 → file descriptor for **stdout**.
- Writes **1 byte** (the character read) to output.

### 3. Loop continues until EOF

- Press **Ctrl+D** (on Linux) to signal **End of File** and terminate input.
- 

## How to Run

```
gcc 6.c -o echo_syscall
```

./echo\_syscall

### Example Interaction

Hello Linux

Hello Linux

- Whatever you type appears immediately on the screen.
- 

### Viva Questions

1. **Q:** What are 0 and 1 in read(0,...) and write(1,...)?

**A:** File descriptors: 0 = stdin, 1 = stdout.

2. **Q:** Difference between read()/write() and scanf()/printf()?

**A:** read/write → low-level system calls, unbuffered.

scanf/printf → high-level C library functions, buffered.

3. **Q:** Why read 1 byte at a time?

**A:** To echo each character immediately; can be changed to read multiple bytes at once.

4. **Q:** How to terminate input?

**A:** Press Ctrl+D (EOF) in Linux terminal.

## Q7. Write a program to copy file1 into file2 (\$cp file1 file2).

### Program Goal

- Copy the contents of **file1.txt** into **file2.txt** (similar to \$ cp file1 file2).
  - Demonstrates use of **low-level system calls**: open(), read(), write(), close().
- 

### Header Files

```
#include <fcntl.h> // for open() and file flags  
#include <stdio.h> // for printf()  
#include <unistd.h> // for read(), write(), close()
```

---

### Constants

```
#define BUF_SIZE 1024
```

- Size of buffer used to read and write data.
  - Reads up to 1024 bytes at a time.
-

## Code (Consolas style)

```
#include <fcntl.h>
#include <stdio.h>
#include <unistd.h>
#define BUF_SIZE 1024

int main() {
    // Open source file in read-only mode
    int fd_src = open("file1.txt", O_RDONLY)

    // Open destination file in write mode, create if not exists, truncate if exists
    int fd_dest = open("file2.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644);

    char buffer[BUF_SIZE];

    int bytesRead;

    // Read from source and write to destination
    while ((bytesRead = read(fd_src, buffer, BUF_SIZE)) > 0) {
        write(fd_dest, buffer, bytesRead);
    }

    printf("Copy Successfully\n");

    // Close both files
    close(fd_src);
    close(fd_dest);

    return 0;
}
```

---

## Step-by-Step Explanation

### 1. Open Files

```
int fd_src = open("file1.txt", O_RDONLY);
int fd_dest = open("file2.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644);
```

- `fd_src` → file descriptor for **source file** (`O_RDONLY` = read only).
- `fd_dest` → file descriptor for **destination file**:
  - `O_WRONLY` → write only
  - `O_CREAT` → create if not exists

- O\_TRUNC → truncate if file exists
  - 0644 → file permission (rw-r--r--)
- 

## 2. Read-Write Loop

```
while ((bytesRead = read(fd_src, buffer, BUF_SIZE)) > 0) {
```

```
    write(fd_dest, buffer, bytesRead);  
}
```

- read() → reads up to BUF\_SIZE bytes from fd\_src into buffer.
  - Returns number of bytes read.
  - write() → writes **exact number of bytes read** to destination.
  - Loop continues until read() returns 0 (EOF).
- 

## 3. Print Confirmation

```
printf("Copy Successfully\n");
```

- Indicates copying finished.
- 

## 4. Close Files

```
close(fd_src);  
close(fd_dest);
```

- Always close file descriptors to **release kernel resources**.
- 

## How to Run

```
gcc 7.c -o filecopy
```

```
./filecopy
```

- Check files:  

```
ls -l file1.txt file2.txt
```
  - Contents of file2.txt will match file1.txt.
- 

## Viva Questions

1. **Q:** Why use O\_TRUNC?

**A:** To empty destination file if it already exists before copying.

2. **Q:** What happens if O\_CREAT is not used and file doesn't exist?

**A:** open() fails; fd\_dest will be -1.

3. **Q:** Why use a buffer?

**A:** Reading/writing byte by byte is slow; buffer allows **block transfer**.

4. **Q:** Difference between read/write and fread/fwrite?

**A:** read/write → low-level, unbuffered system calls.

fread/fwrite → high-level buffered library functions.

5. **Q:** Why do we pass bytesRead to write() instead of BUF\_SIZE?

**A:** Last read may contain fewer bytes than buffer size. Writing exact bytes avoids garbage data.

## **Q8. Write a program to open a file in read only mode, read line by line and display each line as it is read. Close the file when end of file is reached.**

### **Program Goal**

- Open a file in **read-only mode**.
- Read its contents **line by line**.
- Display each line on **standard output**.
- Close the file after reaching **end-of-file (EOF)**.

---

### **Header Files**

```
#include <fcntl.h> // for open() and file flags
```

```
#include <unistd.h> // for read() and write()
```

- open() → to open a file
- read() → to read data from file descriptor
- write() → to write data to stdout
- close() → to close file descriptor

## Code (Consolas Style)

```
#include <fcntl.h>
#include <unistd.h>

int main(int argc, char *argv[]) {
    int fd;
    char c;
    // Check command line arguments
    if (argc != 2) {
        write(1, "this is linux\n", 13);
        return 1;
    }
    // Open file in read-only mode
    fd = open(argv[1], O_RDONLY);
    if (fd == -1) return 1;
    // Read file character by character
    while (read(fd, &c, 1) > 0) {
        write(1, &c, 1); // Display each character
        if (c == '\n') {
            // End of line reached, next iteration reads next line
        }
    }
    // Close file
    close(fd);
    return 0;
}
```

---

## Step-by-Step Explanation

### 1. Command-line Argument

```
if (argc != 2) {
    write(1, "this is linux\n", 13);
    return 1;
```

```
}
```

- Program expects **file name as argument** (argv[1]).
  - If not provided, prints "this is linux" and exits.
- 

## 2. Open File

```
fd = open(argv[1], O_RDONLY);
```

```
if (fd == -1) return 1;
```

- Opens the specified file in **read-only mode**.
  - Returns a **file descriptor** ( $fd \geq 0$ ) on success.
  - Returns -1 on failure (file not found or permission denied).
- 

## 3. Read Character by Character

```
while (read(fd, &c, 1) > 0) {
```

```
    write(1, &c, 1);
```

```
}
```

- `read(fd, &c, 1)` → reads **one character at a time** from file.
- `write(1, &c, 1)` → writes the character to **stdout**.
- Loop continues until `read()` returns 0 (EOF).

Note: The `if(c == '\n')` check is present but not doing anything here. Its purpose could be to **process lines separately** if needed.

---

## 4. Close File

```
close(fd);
```

- Releases the file descriptor and cleans up resources.
- 

## How to Run

```
gcc 8.c -o readline
```

```
./readline myfile.txt
```

## Sample Output

(line-by-line content of myfile.txt displayed)

If no file provided:

./readline

Output:

this is linux

---

### Viva Questions

1. **Q:** Why do we use O\_RDONLY?  
**A:** To open the file only for reading; writing is not allowed.
2. **Q:** Why read 1 byte at a time instead of line at a time?  
**A:** Using read() system call reads raw bytes. To read a full line, you need to check for newline ('\n').
3. **Q:** What does write(1, &c, 1) do?  
**A:** Writes one character to **stdout** (file descriptor 1).
4. **Q:** How is EOF detected?  
**A:** read() returns **0** when end-of-file is reached.
5. **Q:** What happens if file cannot be opened?  
**A:** open() returns -1; program exits.

Q9. Write a program to print the following information about a given file.

- a. inode
- b. number of hard links
- c. uid
- d. gid
- e. size
- f. block size
- g. number of blocks
- h. time of last access
- i. time of last modification
- j. time of last change

### Header Files

```
#include <stdio.h> // printf(), perror()  
  
#include <sys/stat.h> // struct stat, stat()  
  
#include <stdlib.h> // exit()  
  
#include <time.h> // ctime()  
  
#include <unistd.h> // file system related functions
```

---

## Code (Consolas Style)

```
#include <stdio.h>
#include <sys/stat.h>
#include <stdlib.h>
#include <time.h>
#include <unistd.h>

int main() {
    struct stat fileStat;
    const char *filename = "file1.txt";

    // Retrieve file metadata
    if (stat(filename, &fileStat) < 0) {
        perror("stat");
        exit(1);
    }

    // Display file information
    printf("File: %s\n", filename);
    printf("Inode: %ld\n", (long)fileStat.st_ino);
    printf("Number of hard links: %ld\n", (long)fileStat.st_nlink);
    printf("UID: %ld\n", (long)fileStat.st_uid);
    printf("GID: %ld\n", (long)fileStat.st_gid);
    printf("Size: %ld bytes\n", (long)fileStat.st_size);
    printf("Block size: %ld\n", (long)fileStat.st_blksize);
    printf("Number of blocks: %ld\n", (long)fileStat.st_blocks);
    printf("Time of last access: %s", ctime(&fileStat.st_atime));
    printf("Time of last modification: %s", ctime(&fileStat.st_mtime));
    printf("Time of last change: %s", ctime(&fileStat.st_ctime))

    return 0;
}
```

## Step-by-Step Explanation

### 1. stat() system call

```
if (stat(filename, &fileStat) < 0) { perror("stat"); exit(1); }
```

- Retrieves **metadata** of the file and stores in struct stat.
  - Returns 0 on success, -1 on failure.
- 

### 2. File Metadata in struct stat

Field	Description
-------	-------------

st_ino	Inode number (unique file identifier)
--------	---------------------------------------

st_nlink	Number of hard links
----------	----------------------

st_uid	User ID of owner
--------	------------------

st_gid	Group ID of owner
--------	-------------------

st_size	File size in bytes
---------	--------------------

st_blksize	Block size for I/O operations
------------	-------------------------------

st_blocks	Number of allocated blocks
-----------	----------------------------

st_atime	Last access time
----------	------------------

st_mtime	Last modification time
----------	------------------------

st_ctime	Last status change time
----------	-------------------------

### 3. Display times

```
ctime(&fileStat.st_atime)
```

- Converts **epoch time** to human-readable string.
  - Similarly for st\_mtime and st\_ctime.
- 

### 4. Type Casting

```
(long)fileStat.st_ino
```

- Casts fields to long to ensure correct printing on 64-bit systems.
-

## Sample Output

File: file1.txt

Inode: 281474976710963

Number of hard links: 1

UID: 1000

GID: 1000

Size: 12 bytes

Block size: 512

Number of blocks: 0

Time of last access: Fri Sep 5 11:10:28 2025

Time of last modification: Fri Sep 5 11:10:28 2025

Time of last change: Fri Sep 5 11:10:28 2025

---

## Viva Questions

1. **Q:** What is an inode?

**A:** A unique identifier for a file in the filesystem storing metadata (permissions, ownership, timestamps).

2. **Q:** Difference between st\_mtime and st\_ctime?

**A:** st\_mtime → last modification of file content

st\_ctime → last change in file **metadata** (ownership, permissions, link count)

3. **Q:** What is st\_blksize?

**A:** Optimal block size for file I/O operations.

4. **Q:** How many hard links does a newly created file have?

**A:** Usually 1.

5. **Q:** Why cast to (long) when printing?

**A:** To safely print 64-bit fields like inode on all systems.

**Q10. Write a program to open a file with read write mode, write 10 bytes, move the file pointer by 10 bytes (use lseek) and write again 10 bytes.**

**a. check the return value of lseek**

**b. open the file with od and check the empty spaces in between the data.**

## Program Goal

- Open a file in **read-write mode**.
  - Write 10 bytes of data.
  - Move the **file pointer** forward by 10 bytes using lseek().
  - Write another 10 bytes.
  - Observe the **empty spaces (holes)** created in the file.
- 

## Header Files

```
#include <fcntl.h> // open() flags  
#include <unistd.h> // write(), lseek(), close()  
#include <stdio.h> // printf(), perror()
```

---

## Code (Consolas Style)

```
#include <fcntl.h>  
  
#include <unistd.h>  
  
#include <stdio.h>  
  
int main() {  
    int fd  
  
    // Open or create file in read-write mode, truncate if exists  
    fd = open("test10.txt", O_RDWR | O_CREAT | O_TRUNC, 0644);  
  
    // Write first 10 bytes  
    write(fd, "ABCDEFGHIJ", 10)  
  
    // Move file pointer 10 bytes ahead from current position  
    lseek(fd, 10, SEEK_CUR);  
  
    // Write next 10 bytes  
    write(fd, "1234567890", 10);  
  
    close(fd);  
  
    return 0;  
}
```

## Step-by-Step Explanation

### 1. Open File

```
fd = open("test10.txt", O_RDWR | O_CREAT | O_TRUNC, 0644);
```

- O\_RDWR → open for reading and writing
  - O\_CREAT → create file if it does not exist
  - O\_TRUNC → truncate file if it exists
  - 0644 → file permissions (rw-r--r--)
- 

### 2. Write First 10 Bytes

```
write(fd, "ABCDEFGHIJ", 10);
```

- Writes "ABCDEFGHIJ" at **beginning of file** (offset 0–9).
- 

### 3. Move File Pointer

```
lseek(fd, 10, SEEK_CUR);
```

- Moves the **file pointer 10 bytes forward** from the current position.
  - This creates a **hole** of 10 bytes between the first and second write.
  - The return value of lseek() is the **new offset** (useful to verify position).
- 

### 4. Write Next 10 Bytes

```
write(fd, "1234567890", 10);
```

- Writes "1234567890" starting from offset 20–29.
  - The 10 bytes between offsets 10–19 remain as **empty spaces (filled with zeros)**.
- 

### 5. Close File

```
close(fd);
```

- Releases file descriptor.
- 

## Checking the File

Use od (octal dump) to visualize file contents:

```
od -c test10.txt
```

## Output

```
00000000 A B C D E F G H I J \0 \0 \0 \0 \0 \0  
00000020 \0 \0 \0 \0 1 2 3 4 5 6 7 8 9 0  
00000036
```

- Shows **hole (zeros)** between the first and second write.
- 

### Viva Questions

1. **Q:** What does lseek() do?

**A:** Moves the **file pointer** to a specific location in a file.

2. **Q:** What are the arguments of lseek(fd, offset, whence)?

**A:** fd → file descriptor

offset → number of bytes to move

whence → reference point (SEEK\_SET, SEEK\_CUR, SEEK\_END)

3. **Q:** What is a **hole** in a file?

**A:** Unwritten part of a file (contains zeros) created when file pointer is moved beyond current file end.

4. **Q:** Why use O\_TRUNC here?

**A:** Ensures file starts empty each time program runs.

5. **Q:** How to verify empty spaces in the file?

**A:** Use od -c or hexdump -C.

## Q11. Write a program to open a file, duplicate the file descriptor and append the file with both the descriptors and check whether the file is updated properly or not.(1) use dup (2) use dup2 (3) use fcntl

### Program Goal

- Open a file.
  - Duplicate its file descriptor in **three different ways**:
    1. dup()
    2. dup2()
    3. fcntl() with F\_DUPFD.
  - Write to the file using all descriptors.
  - Verify that the file is updated properly (all writes go into the same file).
- 

### Header Files

```
#include <fcntl.h> // open(), fcntl()
```

```
#include <unistd.h> // dup(), dup2(), write(), close()  
#include <stdio.h> // printf(), perror()
```

---

### Code (Consolas style)

```
#include <fcntl.h>  
  
#include <unistd.h>  
  
#include <stdio.h>  
  
int main() {  
  
    int fd1, fd2, fd3, fd4;  
  
    // Open file in append mode  
  
    fd1 = open("test11.txt", O_WRONLY | O_CREAT | O_APPEND, 0644);  
  
    // Duplicate file descriptor using dup()  
  
    fd2 = dup(fd1);  
  
    // Duplicate file descriptor using dup2(), fixed target descriptor (100)  
  
    fd3 = dup2(fd1, 100);  
  
    // Duplicate file descriptor using fcntl()  
  
    fd4 = fcntl(fd1, F_DUPFD, 0);  
  
    // Write to the file using all descriptors  
  
    write(fd1, "Written from fd1\n", 17);  
  
    write(fd2, "Written from fd2 copy from fd1\n", 31);  
  
    write(fd3, "Written from fd3 copy from fd1\n", 31);  
  
    write(fd4, "Written from fd4 copy from fd1\n", 31);  
  
    // Close all descriptors  
  
    close(fd1);  
  
    close(fd2);  
  
    close(fd3);  
  
    close(fd4);  
  
  
    return 0;  
}
```

## Step-by-Step Explanation

### 1. Open the File

```
fd1 = open("test11.txt", O_WRONLY | O_CREAT | O_APPEND, 0644);
```

- Opens (or creates) test11.txt in write-only **append mode**.
  - 0644 → permissions rw-r--r--.
  - First file descriptor returned is fd1.
- 

### 2. Duplicate Using dup()

```
fd2 = dup(fd1);
```

- dup() creates a **copy** of fd1.
  - The **lowest available file descriptor number** is returned.
  - fd1 and fd2 point to the **same open file description** (same file offset, same flags).
- 

### 3. Duplicate Using dup2()

```
fd3 = dup2(fd1, 100);
```

- Creates a duplicate of fd1 in a **specific file descriptor number** (here, 100).
  - If 100 was already open, it would be closed first.
- 

### 4. Duplicate Using fcntl()

```
fd4 = fcntl(fd1, F_DUPFD, 0);
```

- fcntl() with F\_DUPFD duplicates fd1 into the **lowest available descriptor greater than or equal to 0**.
  - Works similar to dup(), but gives you control of the minimum value of new FD.
- 

### 5. Writing to the File

```
write(fd1, "Written from fd1\n", 17);
write(fd2, "Written from fd2 copy from fd1\n", 31);
write(fd3, "Written from fd3 copy from fd1\n", 31);
write(fd4, "Written from fd4 copy from fd1\n", 31);
```

- All descriptors (fd1, fd2, fd3, fd4) point to the **same file**.
- Since the file was opened with **O\_APPEND**, every write goes at the end of the file.

---

## 6. Close File Descriptors

```
close(fd1);  
close(fd2);  
close(fd3);  
close(fd4);
```

- Always good practice to free system resources.
- 

### Sample File Output (test11.txt)

After running the program, the file will contain:

Written from fd1

Written from fd2 copy from fd1

Written from fd3 copy from fd1

Written from fd4 copy from fd1

---

### Viva Questions

1. **Q:** What is the difference between dup() and dup2()?

**A:** dup() → returns lowest available new file descriptor.

dup2() → lets you specify exactly which file descriptor to duplicate into.

2. **Q:** What is the difference between dup() and fcntl(fd, F\_DUPFD, x)?

**A:** dup() → always returns the lowest available FD.

fcntl(F\_DUPFD) → lets you specify a **minimum value** for the new FD.

3. **Q:** Do duplicated file descriptors share file offset?

**A:** Yes. They share the same **open file description** (offset, flags).

4. **Q:** Why did we use O\_APPEND?

**A:** To ensure each write always goes to the end of the file, preventing overwriting.

5. **Q:** What happens if you close fd1 but keep writing with fd2?

**A:** The file remains open via fd2, because file remains open until **all duplicated descriptors are closed**.

## Q12. Write a program to find out the opening mode of a file. Use fcntl.

### 1. File Opening

```
int fd = open("test12.txt", O_RDWR | O_CREAT, 0644);
```

- Opens (or creates if not present) test12.txt with **read & write** mode.
  - Mode 0644 = owner can read/write, group & others can only read.
- 

### 2. Get File Flags

```
int flags = fcntl(fd, F_GETFL);
```

- F\_GETFL → gets **file status flags**.
  - Flags include: access mode (O\_RDONLY, O\_WRONLY, O\_RDWR) + other options like O\_APPEND, O\_NONBLOCK.
- 

### 3. Check Access Mode

```
if ((flags & O_ACCMODE) == O_RDONLY)
    printf("READ ONLY\n");
else if ((flags & O_ACCMODE) == O_WRONLY)
    printf("WRITE ONLY\n");
else if ((flags & O_ACCMODE) == O_RDWR)
    printf("READ & WRITE\n");
```

- O\_ACCMODE mask extracts only the **access mode bits**.
- Program prints whether the file is **read only / write only / read & write**.
- Since we used O\_RDWR, output is:

READ & WRITE

---

#### Sample Output

```
$ gcc 12.c -o q12
$ ./q12
```

READ & WRITE

---

## Viva Questions

1. Q: What does fcntl stand for?  
A: File Control – it provides control operations on file descriptors.
2. Q: What is the purpose of F\_GETFL in fcntl?  
A: It retrieves the file status flags (access mode + other flags).
3. Q: Why do we use the mask O\_ACCMODE?  
A: To extract only the access mode bits (O\_RDONLY, O\_WRONLY, O\_RDWR).
4. Q: What are some common file status flags?  
A: O\_RDONLY, O\_WRONLY, O\_RDWR, O\_APPEND, O\_NONBLOCK, etc.
5. Q: What's the difference between file status flags and file descriptor flags?  
A: Status flags → control how I/O works (e.g., read/write, append).

Descriptor flags → control FD behavior (e.g., FD\_CLOEXEC).

6. Q: Difference between open() and fopen()  
A: open() → system call, low-level, returns int FD.  
fopen() → library function, returns FILE\*, internally uses open().
7. Q: If you open a file with O\_WRONLY | O\_APPEND, what will happen?  
A: You can only write, and every write goes to the end of the file automatically.

## Q13. Write a program to wait for STDIN for 10 seconds using select. Print whether data is available within 10 seconds or not.

### Code Explanation

```
#include <stdio.h>
#include <unistd.h>
#include <sys/select.h>
```

- Includes standard I/O, UNIX system calls (read), and the **select()** system call library.

---

```
int main(){
    fd_set set;
    struct timeval timeout = {10, 0};
```

- fd\_set set; → A special data structure used by select() to monitor file descriptors (FDs).
- struct timeval timeout = {10, 0}; → Sets the timeout to **10 seconds** (10 sec, 0 microseconds).

---

```
    FD_ZERO(&set);
    FD_SET(0, &set);
```

- FD\_ZERO(&set) → Clears the set (removes all file descriptors).
  - FD\_SET(0, &set) → Adds **file descriptor 0** (which is STDIN) into the set.  
👉 Means we are monitoring **keyboard input**.
- 

```
printf("Enter input within 10 seconds:\n");
if(select(1, &set, NULL, NULL, &timeout) > 0) {
```

- select(nfds, &readfds, &writefds, &exceptfds, &timeout) → Waits until one of the given FDs is **ready for I/O** or until timeout expires.
    - Here:
      - nfds = 1 → Highest FD to check is 0 (stdin), so we pass 1.
      - &set → monitor read events on stdin.
      - NULL, NULL → not checking write or exception conditions.
      - &timeout → wait at most 10 seconds.
  - If return value > 0 → Some FD (here stdin) is ready for reading.
- 

```
char buf[100];
read(0, buf, sizeof(buf));
printf("You typed: %s", buf);
```

- If input is available → Read up to 100 characters from stdin into buf.
  - Print back the user's input.
- 

```
} else {
    printf("No input in 10 seconds\n");
}
return 0;
}
```

- If timeout expires → Print message "No input in 10 seconds".
- 

## ✓ Program Behavior

- If the user types something within **10 seconds**, it prints the input.
- If no input, it prints **timeout message**.

## Viva Questions

1. **Q:** What is the purpose of select() system call?

**A:** select() allows monitoring multiple file descriptors simultaneously to see if they are ready for reading, writing, or exceptional conditions.

2. **Q:** What is the role of struct timeval?

**A:** It specifies the timeout value (in seconds and microseconds) for how long select() should wait.

4. **Q:** What happens if we pass NULL as timeout in select()?

**A:** select() will wait **indefinitely** until at least one FD becomes ready.

5. **Q:** What is the difference between FD\_ZERO, FD\_SET, and FD\_ISSET?

**A:** FD\_ZERO → Clears all FDs from the set.

FD\_SET(fd, &set) → Adds fd into the set.

FD\_ISSET(fd, &set) → Checks if fd is in the set after select() returns.

6. **Q:** Why is the first argument to select() 1 here?

**A:** It should be the highest FD in any set + 1. Since only stdin (0) is monitored, the highest is 0, so 1 is passed.

7. **Q:** What are the possible return values of select()?

**A:** > 0 → Number of ready FDs.

0 → Timeout occurred.

-1 → Error.

8. **Q:** How is select() different from poll() or epoll()?

**A:** select() → Limited by FD\_SETSIZE (typically 1024), less scalable.

poll() → Uses an array of FDs, no fixed size limit.

epoll() → More efficient for large number of FDs, available in Linux only.

9. **Q:** What happens if input arrives after the 10-second timeout?

**A:** The program would already print "No input in 10 seconds" and exit, so late input is ignored.

## Code Explanation

```
#include <stdio.h>
```

```
#include <sys/stat.h>
```

- stdio.h → Standard I/O.
- sys/stat.h → Declares the stat structure and macros to test file types.

```
int main(int argc, char *argv[]) {
```

```
struct stat st;
```

- argc, argv → Used to take filename as command-line argument.
  - struct stat st; → Holds file metadata (permissions, type, size, etc.).
- 

```
if (argc != 2){  
    printf("Usage: ./task14 <filename>\n");  
    return 1;  
}
```

- Ensures that exactly **one filename** is passed as argument.
  - Example: ./task14 myfile.txt.
- 

```
if (stat(argv[1], &st) == -1){  
    perror("stat");  
    return 1;  
}
```

- stat(path, &st) → Fills st with file info.
  - If it fails (file not found, no permission, etc.), print error with perror.
- 

```
if (S_ISREG(st.st_mode))  
    printf("Regular file\n");  
else if (S_ISDIR(st.st_mode))  
    printf("Directory\n");  
else if (S_ISLNK(st.st_mode))  
    printf("Symbolic link\n");  
else if (S_ISCHR(st.st_mode))  
    printf("Character device\n");  
else if (S_ISBLK(st.st_mode))  
    printf("Block device\n");  
else if (S_ISFIFO(st.st_mode))  
    printf("FIFO (named pipe)\n");  
else if (S_ISSOCK(st.st_mode))
```

```
printf("Socket\n");
else
    printf("Unknown file type\n");
```

- These macros check st.st\_mode bits to determine file type:
  - S\_ISREG → Regular file
  - S\_ISDIR → Directory
  - S\_ISLNK → Symbolic link
  - S\_ISCHR → Character device (e.g., /dev/tty)
  - S\_ISBLK → Block device (e.g., /dev/sda)
  - S\_ISFIFO → FIFO (named pipe)
  - S\_ISSOCK → Socket

### Program Behavior

- If you run: ./task14 myfile.txt → Regular file
- If you run: ./task14 /etc → Directory
- If you run: ./task14 /dev/null → Character device

### Viva Questions

1. **Q:** What is the purpose of stat()?  
**A:** stat() retrieves metadata about a file (size, type, permissions, timestamps, etc.).
2. **Q:** What happens if the file does not exist?  
**A:** stat() returns -1, and errno is set (e.g., ENOENT for "No such file").
3. **Q:** What does S\_ISREG(st.st\_mode) check?  
**A:** Whether the file is a **regular file** (like .txt, .c, etc.).

## Q15. Write a program to display the environmental variable of the user (use environ).

### Code Explanation

```
#include <stdio.h>
```

- Standard I/O functions like printf.

```
extern char **environ;
```

- environ is a **global variable** provided by the C runtime.

- It is an array of C strings (`char*`), each string representing an environment variable in the form:
- `KEY=VALUE`

Example:

`PATH=/usr/bin:/bin`

`HOME=/home/jils`

`USER=jils`

---

```
int main() {
    int i = 0;
    while (environ[i] != NULL) {
        printf("%s\n", environ[i]);
        i++;
    }
    return 0;
}
```

- Starts at index 0.
- Loops until it finds `NULL` (end of array).
- Prints each environment variable string.

#### Program Output Example

`PATH=/usr/bin:/bin:/usr/local/bin`

`HOME=/home/jils`

`USER=jils`

`SHELL=/bin/bash`

`LANG=en_US.UTF-8`

...

---

#### Viva Questions

1. **Q:** What does `environ` contain?

**A:** It contains the **environment variables** of the process, stored as an array of strings in `KEY=VALUE` format.

2. **Q:** What is an environment variable?

**A:** It is a key-value pair that affects how processes run. Example: PATH defines where executables are searched.

3. **Q:** How do you access environment variables in C besides environ?

**A:** Using getenv("VAR\_NAME").

4. **Q:** Difference between argv and environ?

**A:**

- o argv → Command-line arguments given by the user at program start.
- o environ → Inherited environment variables from the parent process (usually the shell).

5. **Q:** How can you set or modify an environment variable inside C?

**A:** Using functions like setenv(), putenv(), or unsetenv().

6. **Q:** Where are environment variables stored in memory?

**A:** Typically stored in the process's **user space** (part of the process's memory image set up by the OS when the program starts).

7. **Q:** What happens to environment variables when a process creates a child process?

**A:** They are **inherited** by the child (copied).

8. **Q:** Are environment variables global across all processes?

**A:** No. They are **per-process**, but inherited from parent to child. Changing them in one process does not affect others.

9. **Q:** Can we remove extern char \*\*environ; and still get environment variables?

**A:** Yes, if we declare main as int main(int argc, char \*argv[], char \*envp[]), then envp provides the same environment strings.

## Q16. Write a program to perform mandatory locking. a. Implement write lock b. Implement read lock

16a.c – Write Lock

Purpose

This program demonstrates mandatory file locking in Linux using a write lock. A write lock ensures that no other process can read or write to the file until the lock is released.

Step-by-step Explanation

1. Header files

```
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <stdlib.h>
```

- stdio.h → for printf and scanf

- unistd.h → for write, read, close
  - fcntl.h → for fcntl and file locking (struct flock)
  - stdlib.h → for exit
- 

## 2. Open the file

```
fd = open("text16a.txt", O_RDWR | O_CREAT, 0666);
```

- Opens text16a.txt in read-write mode.
  - Creates the file if it doesn't exist (O\_CREAT).
  - Permission 0666 → read & write for all users.
  - Returns file descriptor fd.
- 

## 3. Set up the lock

```
lock.l_type = F_WRLCK; // Write lock  
  
lock.l_whence = SEEK_SET; // Start of file  
  
lock.l_start = 0; // Start offset  
  
lock.l_len = 0; // Lock entire file  
  
lock.l_pid = getpid(); // Current process ID
```

- F\_WRLCK → write lock
  - SEEK\_SET → lock starts from beginning
  - l\_len = 0 → lock whole file
  - l\_pid → needed for identifying the process holding the lock
- 

## 4. Acquire the lock

```
fcntl(fd, F_SETLKW, &lock);
```

- F\_SETLKW → wait until lock is available (blocking call)
  - Locks the file for writing.
- 

## 5. Write data

```
scanf("%s", buffer);  
  
write(fd, buffer, sizeof(buffer));
```

- User inputs text and writes it to the file.

- Only one process can write at a time due to the write lock.
- 

## 6. Release the lock

```
lock.l_type = F_UNLCK;  
fcntl(fd, F_SETLK, &lock);
```

- F\_UNLCK → unlocks the file.
  - F\_SETLK → non-blocking unlock.
- 

## 7. Close the file

```
close(fd);
```

### Key Points

- Write lock prevents other processes from reading or writing.
  - Blocking vs non-blocking locks: F\_SETLKW waits; F\_SETLK fails if lock unavailable.
- 

## 2 16b.c – Read Lock

### Purpose

This program demonstrates mandatory file locking in Linux using a read lock. A read lock allows multiple processes to read, but no process can write while the read lock is active.

---

### Step-by-step Explanation

#### 1. Open file

```
fd = open("text16a.txt", O_RDONLY);
```

- Open the file in read-only mode.
- 

#### 2. Set up the lock

```
lock.l_type = F_RDLCK; // Read lock  
lock.l_whence = SEEK_SET;  
lock.l_start = 0;  
lock.l_len = 0;  
lock.l_pid = getpid();
```

- F\_RDLCK → read lock

- Multiple read locks can coexist.
- 

### 3. Acquire lock

```
fcntl(fd, F_SETLKW, &lock);
```

- Blocks until write locks are released.
  - Other processes can also acquire read locks simultaneously.
- 

### 4. Read file

```
read(fd, buffer, sizeof(buffer));
```

- Reads the file content into buffer.
- 

### 5. Release lock

```
lock.l_type = F_UNLCK;
```

```
fcntl(fd, F_SETLK, &lock);
```

---

### 6. Close file

```
close(fd);
```

#### Key Points

- Read locks allow concurrent reading but block writing.
- Write locks are exclusive.

**Q17. Write a program to simulate online ticket reservation. Implement write lock Write a program to open a file, store a ticket number and exit. Write a separate program, to open the file, implement write lock, read the ticket number, increment the number and print the new ticket number then close the file.**

17a.c – Initialize Ticket Number

Purpose

- This program creates a file db and stores an initial ticket number (10).
- Acts as the database initialization for ticket reservation.

Step-by-Step Explanation

#### 1. Header files

```
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <stdlib.h>
```

## 2. Structure

```
struct {
    int ticket_no;
} db;
```

- db holds the ticket number.

## 3. Initialize ticket number

```
db.ticket_no = 10;
```

## 4. Open or create file

```
fd = open("db", O_CREAT | O_RDWR, 0766);
```

- Opens db for read/write.
- Creates it if not present.
- Permissions 0766 → owner can read/write/execute, others read/write.

## 5. Write ticket number to file

```
write(fd, &db, sizeof(db));
```

- Stores the structure in the file.

## 6. Print confirmation

```
printf("Initial ticket number stored: %d\n", db.ticket_no);
```

## 7. Close file

```
close(fd);
```

- Key Point: This program does not use locking, because it just initializes the ticket number.

---

## 2 17b.c – Ticket Reservation with Write Lock

### Purpose

- Simulates online ticket booking using critical section protection.
- Ensures only one process can update ticket number at a time.

---

### Step-by-Step Explanation

1. Open file

```
fd = open("db", O_RDWR);
```

- Open the file for reading and writing.

2. Set up write lock

```
lock.l_type = F_WRLCK; // Write lock  
lock.l_whence = SEEK_SET;  
lock.l_start = 0;  
lock.l_len = 0;      // Entire file  
lock.l_pid = getpid();
```

- Prepares a write lock on the whole file.

3. Acquire lock

```
fcntl(fd, F_SETLKW, &lock);
```

- Blocks if another process has a write lock.
- Only one process can enter critical section.

---

4. Move file pointer to start

```
lseek(fd, 0, SEEK_SET);
```

- Ensures reading starts from the beginning of file.

5. Read current ticket number

```
read(fd, &db, sizeof(db));  
printf("Current ticket Number: %d\n", db.ticket_no);
```

6. Increment ticket number

```
db.ticket_no++;  
lseek(fd, 0, SEEK_SET);  
write(fd, &db, sizeof(db));  
printf("New ticket Number: %d\n", db.ticket_no);
```

- Updates the ticket number and writes back to the file.

7. Release lock

```
lock.l_type = F_UNLCK;  
fcntl(fd, F_SETLK, &lock);
```

- Unlocks the file so other processes can enter the critical section.

## 8. Close file

`close(fd);`

- Key Point: The write lock ensures mutual exclusion, preventing race conditions if multiple processes try to book tickets simultaneously.
- 

## 3 Important Concepts

Concept	Explanation
struct flock	Data structure for file locking.
F_WRLCK	Write lock (exclusive).
F_SETLKW	Blocking lock acquisition.

`lseek(fd, 0, SEEK_SET)` Move file pointer to start before reading/writing.

Critical Section (CS) Code that must be executed by only one process at a time.

---

## 4 Expected Viva Questions with Answers

**Q1: What is the purpose of 17a.c?**

**A: It initializes the ticket number in a file db before any reservation.**

**Q2: What does 17b.c do?**

**A: It reads the ticket number, increments it (simulating ticket booking), and writes back using a write lock to prevent concurrent access.**

**Q3: Why do we use F\_WRLCK?**

**A: To prevent other processes from reading or writing while updating the ticket number (mutual exclusion).**

**Q4: What is the purpose of lseek(fd, 0, SEEK\_SET)?**

**A: To move the file pointer to the beginning before reading or writing.**

**Q5: Why do we use F\_SETLKW instead of F\_SETLK?**

**A: F\_SETLKW blocks the process until the lock is available. Ensures no two processes enter the critical section simultaneously.**

**Q6: What happens if two processes run 17b.c simultaneously?**

**A: Only one enters the critical section. The second waits until the lock is released. This prevents race conditions.**

**Q7: What happens if we don't use a lock?**

**A: Multiple processes could read the same ticket number simultaneously, causing duplicate ticket numbers (race condition).**

**Q8: What is a critical section in this program?**

**A: The part where the ticket number is read, incremented, and written back. Only one process should execute it at a time.**

**Q9: Can multiple processes read the ticket number simultaneously?**

**A: If we use read lock (F\_RDLCK) instead of write lock, yes. But bookijyang requires write lock to avoid duplication.**

**Q10: What is the role of getpid() in l\_pid?**

**A: Identifies which process owns the lock.**

**Q18. Write a program to perform Record locking. (a) Implement write lock (b) Implement read lock Create three records in a file. Whenever you access a particular record, first lock it then modify/access to avoid race condition.**

18a.c – Create Train Records

Purpose

- Creates a file record.txt with 3 train records.
- Each train has:
  - train\_num → train ID (1, 2, 3)
  - ticket\_count → number of tickets booked (initially 0).

Explanation

1. Define a structure:

```
struct {  
    int train_num;  
    int ticket_count;  
} db[3];
```

- Array of 3 trains.

2. Initialize trains:

```
for(int i=0;i<3;i++){  
    db[i].train_num=i+1;  
    db[i].ticket_count=0;  
}
```

3. Create/open file record.txt:

```
fd=open("record.txt",O_CREAT|O_RDWR,0666);
```

4. Write all 3 records to file:

```
write(fd,db,sizeof(db));
```

Result: A file with 3 train records is created.

---

- ◆ 18b.c – Book Ticket (Write Lock)

#### Purpose

- Allows user to select a train and book a ticket.
- Uses a write lock so only one process can update a train record at a time.

#### Explanation

1. Open file:

```
fd=open("record.txt",O_RDWR);
```

2. Ask user to choose train (1–3):

```
scanf("%d",&input);
```

3. Define lock for selected train:

```
lock.l_type=F_WRLCK; // Write lock  
lock.l_whence=SEEK_SET;  
lock.l_start=(input-1)*sizeof(db); // start offset  
lock.l_len[sizeof(db)]; // lock only that train's record
```

 This ensures only the chosen train record is locked, not the whole file.

4. Move pointer and read train record:

```
lseek(fd,(input-1)*sizeof(db),SEEK_SET);  
read(fd,&db,sizeof(db));
```

5. Acquire lock:

```
fctl(fd,F_SETLKW,&lock);
```

6. Update ticket count:

```
db.ticket_count++;  
lseek(fd,(input-1)*sizeof(db),SEEK_SET);  
write(fd,&db,sizeof(db));
```

7. Release lock after booking.

Result: Safely increments ticket count for selected train.

---

- ◆ 18c.c – View Ticket Count (Read Lock)

#### Purpose

- Allows user to view current ticket count of a train.
- Uses read lock → multiple people can read simultaneously, but no one can write while reading.

#### Explanation

1. Open file in read-only mode:

```
fd=open("record.txt",O_RDONLY);
```

2. Ask for train number.

3. Define read lock:

```
lock.l_type=F_RDLCK; // Read lock  
lock.l_start=(input-1)*sizeof(db);  
lock.l_len=sizeof(db);
```

4. Acquire lock:

```
fctl(fd,F_SETLK,&lock);
```

5. Read the ticket count and display:

```
read(fd,&db,sizeof(db));  
printf("Train %d - Current ticket count: %d\n", db.train_num, db.ticket_count);
```

6. Unlock and exit.

- Result: Safely reads ticket count without disturbing booking.
- 

- ◆ Summary of Three Codes

- 18a.c → Initializes trains (database setup).
  - 18b.c → Simulates ticket booking (with write lock).
  - 18c.c → Simulates checking ticket count (with read lock).
- 

#### Expected Viva Q&A

Q1: What is the purpose of 18a.c?

A: To create record.txt and initialize 3 train records with ticket\_count=0.

Q2: Why do we need locks in 18b.c and 18c.c?

A: To avoid race conditions. Without locks, two people booking at the same time might overwrite each other's changes.

Q3: What is the difference between write lock and read lock here?

- Write lock (F\_WRLCK) → Only 1 process can book a ticket at a time (exclusive).
- Read lock (F\_RDLCK) → Many processes can view ticket count at the same time, but booking (write) is blocked.

Q4: Why do we use lseek in these programs?

A: To move the file pointer to the correct train record before reading/writing.

Q5: What does this mean?

```
lock.l_start=(input-1)*sizeof(db);
```

```
lock.l_len=sizeof(db);
```

A: It means lock only one train's record, not the whole file.

Q6: What will happen if two people run 18b.c for the same train at the same time?

A: One will enter critical section, the other will wait (because of F\_SETLKW).

Q7: What happens if one process runs 18b.c (write lock) and another runs 18c.c (read lock) for the same train?

A: The read process will wait until the write lock is released.

Q8: What happens if two processes run 18c.c at the same time?

A: Both can read simultaneously since read locks are shared.

Q9: Why is sizeof(db) used in write and read?

A: To ensure the entire structure (train record) is read/written, not just part of it.

Q10: What is the role of getpid() in lock.l\_pid?

A: To record which process holds the lock (useful for debugging).

## **Q19. Write a program to find out time taken to execute getpid system call. Use time stamp counter.**

### **Explanation of the Code**

```
#include <stdio.h>
#include <unistd.h>
#include <x86intrin.h>
int main(){
    unsigned long long start,end;
    start = _rdtsc(); // Read CPU cycle counter before getpid()
```

```

getpid();      // System call: get process ID
end = _rdtsc(); // Read CPU cycle counter after getpid()
printf("Total Cycles: %llu\n", end - start);
return 0;

```

### Step-by-step:

1. `_rdtsc()`:
    - o Reads the **Time Stamp Counter** (TSC), which counts the number of CPU cycles since reset.
    - o Returns a 64-bit unsigned integer.
  2. `getpid()`:
    - o A **system call** that returns the current process ID.
    - o Since it's a **kernel interaction**, it takes more cycles than normal instructions.
  3. **Cycle Calculation:**
    - o `end - start` gives the number of CPU cycles taken by the `getpid()` system call.
- 

### ◆ Sample Output

If you run it, you may see output like:

Total Cycles: 1200

(The number of cycles varies depending on CPU, kernel version, caching, and system load.)

---

### ◆ Expected Viva Questions (with Answers)

#### Conceptual

1. **Q: What does `getpid()` do?**

A: It returns the process ID of the calling process.
2. **Q: What is `_rdtsc()`?**

A: `_rdtsc()` is an intrinsic instruction (based on the RDTSC assembly instruction) that reads the CPU's Time Stamp Counter, which counts the number of cycles since the CPU was powered on/reset.
3. **Q: Why use TSC instead of `time()` or `clock()`?**

A: TSC gives **high-resolution timing in CPU cycles**, while `time()/clock()` give coarse-grained measurements in seconds or microseconds.

**4. Q: Is getpid() a system call or a library function?**

A: It is a **system call**, but usually accessed through a wrapper function in the C standard library (glibc).

**5. Q: Why does a system call take more time than a normal function call?**

A: Because it involves **switching from user mode to kernel mode** and then back to user mode, which adds overhead.

**Q20.**

**Objective**

The program is about:

1. Finding the **priority (nice value)** of the currently running process.
  2. Modifying it by **increasing the nice value** (which effectively lowers the priority).
- 

◆ **Code Explanation**

```
#include <stdio.h>
#include <sys/time.h>
#include <sys/resource.h>
#include <unistd.h>

int main() {
    int priority;

    // Get current priority (nice value) of this process
    priority = getpriority(PRIO_PROCESS, 0);
    printf("Old priority: %d\n", priority);

    // Try to increase nice value by 5
    if (setpriority(PRIO_PROCESS, 0, priority + 5) == -1) {
        perror("setpriority");
    }
    else {
        printf("New priority to: %d\n", priority + 5);
    }
}
```

```
}

return 0;
}
```

---

## Step-by-Step Flow

### 1. **getpriority(PRIO\_PROCESS, 0)**

- PRIO\_PROCESS means we are querying based on **process ID**.
- 0 means **current process**.
- Returns the current **nice value** of the process.
- Default nice value = **0**.

👉 Nice value range: **-20 (highest priority)** to **19 (lowest priority)**.

---

### 2. **setpriority(PRIO\_PROCESS, 0, priority + 5)**

- Tries to set the nice value to priority + 5.
- If old nice = 0, new nice = 5.
- **Higher nice → lower scheduling priority.**
- Only root (superuser) can decrease nice value (make it higher priority), but normal users can increase nice value (make it lower priority).

### 3. **Error Handling**

- If setpriority fails (e.g., due to permissions), perror("setpriority") prints the error message.
- Example: If a non-root user tries to reduce the nice value below 0 → Permission denied.

#### ◆ **Example Run**

\$ ./a.out

Old priority: 0

New priority to: 5

If run as a normal user, you can **increase nice**, but cannot decrease it.

## Viva Questions

### 1. Q: What is a process priority in Linux?

A: It determines how often a process gets CPU time. Higher priority = more CPU time.

### 2. Q: What is a nice value?

A: It's a user-space abstraction for process priority, ranging from -20 (highest) to 19 (lowest). Default is 0.

### 3. Q: Difference between priority and nice value?

A: Nice value is a user-defined hint; the kernel converts it internally to a scheduling priority.

### 4. Q: What does increasing the nice value mean?

A: It lowers the priority → process runs less frequently compared to others.

### 5. Q: Why can't normal users set negative nice values?

A: Because giving high priority (-nice) can hog CPU and affect system stability. Only root can do that.

### 6. Q: Alternative way to set nice value from shell?

A: Use nice or renice commands. Example:

- nice -n 10 ./a.out (start with nice value 10)
- renice 5 -p 1234 (change process 1234 to nice value 5).

## Q21. Write a program, call fork and print the parent and child process id.

### Code Explanation

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(){
    pid_t pid;
    pid=fork(); // create a new process
    if(pid<0){ // fork failed
        printf("error");
        return 1;
    }
    else if(pid==0){ // Child process (fork() returns 0 to child)
        printf("Child Process: PID = %d, Parent PID = %d\n", getpid(), getppid());
    }
}
```

```

}

else { // Parent process (fork() returns child's PID)

    printf("Parent Process: PID = %d, Child PID = %d\n", getpid(), pid);

}

return 0;

}

```

## Step-by-Step Working

### 1. fork()

- o Creates a new process (child) that is a duplicate of the parent.
- o Returns:
  - >0 → Child's PID (in parent process)
  - 0 → In child process
  - <0 → Error (fork failed)

### 2. Parent Process Block (pid > 0)

- o Executes the else block.
- o Prints its own PID (getpid()) and the PID of the child (pid).

### 3. Child Process Block (pid == 0)

- o Executes the else if block.
- o Prints its own PID (getpid()) and its parent's PID (getppid()).

## Viva Questions

### 1. Q: What does fork() do?

A: Creates a new process (child) which is almost identical to the parent.

### 2. Q: How does the return value of fork() differ between parent and child?

A: 0 in child, child's PID in parent, -1 on error.

### 3. Q: What are getpid() and getppid()?

A: getpid() → current process ID, getppid() → parent process ID.

### 4. Q: Does child process get a copy of parent's memory?

A: Yes, but with **copy-on-write** mechanism (they share pages until modified).

### 5. Q: Can the parent terminate before the child? What happens then?

A: Yes, then child becomes a **zombie** until parent collects its exit status, or gets re-parented to init/systemd.

**6. Q: How do you prevent zombie processes?**

A: Use wait() or waitpid() in the parent.

**Q22. Write a program, open a file, call fork, and then write to the file by both the child as well as the parent processes. Check output of the file.**

**Code Explanation**

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <stdlib.h>

int main() {
    int fd;
    pid_t pid;

    // Open file in write-only + append mode, create if not exists
    fd = open("output22.txt", O_CREAT | O_WRONLY | O_APPEND, 0644);
    pid = fork(); // create child process

    if (pid == 0) { // Child process
        char child_msg[] = "Message from Child process\n";
        write(fd, child_msg, sizeof(child_msg) - 1);
        printf("Child (PID=%d) wrote to file.\n", getpid());
    }
    else { // Parent process
        char parent_msg[] = "Message from Parent process\n";
    }
}
```

```

        write(fd, parent_msg, sizeof(parent_msg) - 1);

        printf("Parent (PID=%d) wrote to file.\n", getpid());

    }

close(fd); // both processes close their copy of fd

return 0;
}

```

#### ◆ Important Concepts

##### 1. **open()**

- Opens/creates the file output22.txt.
- Flags:
  - O\_CREAT: create file if it doesn't exist.
  - O\_WRONLY: write-only mode.
  - O\_APPEND: append new data to the end of the file.
- 0644: file permissions (rw-r--r--).

##### 2. **fork()**

- Creates a child process.
- Both parent and child share the **same file descriptor** (but in separate process tables).

##### 3. **write()**

- Both processes write to the same file.
- Since file is opened in O\_APPEND, each write goes at the **end of file**.

##### 4. **close()**

- Each process closes its file descriptor after writing.

#### ◆ Example Run

\$ ./a.out

Parent (PID=2345) wrote to file.

Child (PID=2346) wrote to file.

**output22.txt content might look like:**

Message from Parent process

## Message from Child process

⚠ **Order is not guaranteed** — sometimes the child may write first, sometimes the parent, depending on scheduling.

### Viva Questions

#### 1. Q: What happens to file descriptors after a fork()?

A: They are **copied** into the child process. Both processes share the same open file description (file offset, flags, etc.).

#### 2. Q: Why use O\_APPEND?

A: To ensure both parent and child write at the end of the file safely, without overwriting each other's data.

#### 3. Q: Can parent and child write simultaneously to the file?

A: Yes, but writes of less than PIPE\_BUF (usually 4096 bytes) are **atomic** in Linux, so they won't interleave characters. However, the order is nondeterministic.

#### 4. Q: What if we removed O\_APPEND?

A: Then both parent and child would share the same file offset → possible race conditions → overlapping writes.

#### 5. Q: What's the difference between open() and fopen()?

A: open() is a low-level system call (returns file descriptor), fopen() is a C library call (returns FILE\* and provides buffering).

#### 6. Q: How can you ensure parent writes after child?

A: Use wait(NULL) in parent, so child finishes first.

## Q23. Write a program to create a Zombie state of the running program.

### Code Walkthrough

```
#include <stdio.h>
#include <sys/wait.h>
#include <unistd.h>

int main() {
    if (!fork()) {
        // Child process (fork returns 0)
        printf("In child - PID = %d\n", getpid());
    }
    else {
```

```

// Parent process

printf("This is parent process %d\n", getpid());

sleep(30); // Parent sleeps for 30 seconds

wait(0); // After sleep, parent reaps the child

}

}

```

---

#### ◆ Step-by-Step Execution

##### 1. fork()

- Creates a child process.
- Child executes the first block (printf("In child...")).
- Parent executes the second block.

##### 2. Child process

- Prints its PID.
- Then immediately exits → becomes a **zombie** (since parent hasn't wait(ed yet)).

##### 3. Parent process

- Prints its PID.
  - Sleeps for **30 seconds**. During this time:
    - Child is already finished.
    - Entry of child remains in process table as **Zombie (Z)**.
  - After sleep, parent calls wait(0) → collects child's exit status → zombie is removed.
- 

#### ◆ How to See Zombie in Action

1. Compile & run:
2. gcc zombie\_demo.c -o zombie\_demo
3. ./zombie\_demo
4. While the parent is sleeping (30s window), open another terminal and run:
5. ps -l

You will see the child with **STAT = Z** (zombie).

Example:

F S	UID	PID	PPID	C PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
-----	-----	-----	------	-------	----	------	----	-------	-----	------	-----

```
OZ 1000 12346 12345 0 80 0 - 0 exit pts/0 00:00:00 zombie_demo <defunct>
```

6. After 30 seconds, parent calls wait(0) → zombie disappears from the process table.

## Key Points for Viva

### 1. What's happening here?

- Child exits quickly → becomes zombie.
- Parent is still sleeping → not calling wait() yet.
- Hence zombie exists for ~30s.

### 2. Why does zombie disappear after 30s?

- Because parent finally calls wait(0) → kernel removes child's entry.

### 3. Why is wait() important?

- To reap child processes and prevent zombies.

### 4. What if parent never calls wait()?

- Zombie stays until parent exits.
- When parent exits, child gets re-parented to init (or systemd), which reaps it automatically.

## Q24. Write a program to create an orphan process.

### Code Walkthrough

```
#include <stdio.h>
#include <sys/wait.h>
#include <unistd.h>
#include <stdlib.h>

int main() {
    if (!fork()) { // Child process
        printf("In child - Before orphan PPID = %d\n", getppid());
        sleep(1); // Wait so parent exits first
        printf("In child - After orphan PPID = %d\n", getppid());
    }
    else { // Parent process
        printf("This is parent process %d\n", getpid());
    }
}
```

```
    exit(0); // Parent exits immediately  
}  
}
```

---

#### ◆ Step-by-Step Execution

##### 1. Parent creates child using fork().

- Child enters first block (`!fork()` is true → returns 0).
- Parent enters the else block.

##### 2. Child (before orphaning)

- Prints its parent PID with `getppid()` → this will be the parent's PID.

##### 3. Parent exits

- Calls `exit(0)` right away.
- Parent process is gone → child becomes **orphan**.

##### 4. Child (after orphaning)

- After `sleep(1)`, child wakes up and calls `getppid()` again.
  - Now parent PID is 1 (on Linux, it's init or systemd), meaning child got **re-parented**.
- 

#### ◆ Example Output

This is parent process 2345

In child - Before orphan PPID = 2345

In child - After orphan PPID = 1

---

#### Key Viva Questions

##### 1. Q: What is an orphan process?

A: A child process whose parent has terminated. The orphan is re-parented to init (PID 1), which takes care of it.

##### 2. Q: What happens to an orphan process?

A: It continues execution normally. It's adopted by init/systemd.

##### 3. Q: Difference between Zombie and Orphan?

- **Zombie:** Child finished, parent still alive but hasn't collected its status → entry remains in process table (STAT=Z).
- **Orphan:** Parent died first, child is still running → re-parented to init (or systemd).

#### 4. Q: How can you verify an orphan process?

A: Run ps -ef or ps -l → check child's PPID. It will be 1 if it's orphaned.

### Q25. Write a program to create three child processes. The parent should wait for a particular child (use waitpid system call).

#### Included Headers

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/wait.h>
#include <unistd.h>
```

- stdio.h → for printf()
- stdlib.h → for exit()
- sys/wait.h → for waitpid() and macros like WIFEXITED(), WEXITSTATUS()
- unistd.h → for fork(), getpid(), sleep()

---

#### Main Function

```
int main() {
    pid_t cpid[3];
    for (int i = 0; i < 3; i++) {
        cpid[i] = fork();
        if (cpid[i] == 0) {
            printf("Child %d (PID: %d) created\n", i + 1, getpid());
            sleep(2);
            exit(0);
        }
    }
}
```

- pid\_t cpid[3]; → stores PIDs of the 3 child processes.

- for loop runs 3 times to create 3 children using fork().
  - fork() returns:
    - 0 → if we are in the **child process**
    - 0 → **PID of child** in the parent process
  - Inside child (cpid[i] == 0):
    - Prints a message
    - sleep(2) → delay to simulate work
    - exit(0) → terminate child with status 0
  - After fork(), the **parent** continues the loop without executing the if block.
- 

## Parent Waiting for Child 2

```
printf("Parent PID: %d waiting for Child 2 (PID: %d)\n", getpid(), cpid[1]);\n\nint status;\n\npid_t wpid = waitpid(cpid[1], &status, 0);\n\nif (WIFEXITED(status)) {\n    printf("Child 2 (PID: %d) exited with status %d\n", wpid, WEXITSTATUS(status));\n}
```

- getpid() → gets the **parent's PID**
  - cpid[1] → PID of the **second child** (arrays are 0-indexed)
  - waitpid(cpid[1], &status, 0) → waits **specifically for Child 2** to terminate.
  - WIFEXITED(status) → checks if the child exited normally.
  - WEXITSTATUS(status) → returns the **exit code** of child.
- 

## Return Statement

```
return 0;
```

- Indicates the program ended successfully.
- 

## Output

Child 1 (PID: 390) created

Parent PID: 389 waiting for Child 2 (PID: 391)

Child 2 (PID: 391) created

Child 3 (PID: 392) created

Child 2 (PID: 391) exited with status 0

- Shows that **Child 2 finished** and parent waited specifically for it.
  - Order of child creation may vary slightly due to process scheduling.
- 

## Key Points

1. fork() → creates a new process.
2. waitpid() → allows parent to wait for a **specific child**.
3. WIFEXITED() and WEXITSTATUS() → used to check **exit status**.
4. Each child should call exit() to terminate properly.
5. Parent continues execution for remaining children unless explicitly waited.

## Viva Questions

1. **Q:** How many processes are created in this program?  
**A:** 4 processes in total – 1 parent and 3 children.
2. **Q:** What is waitpid()?  
**A:** A system call that allows the parent to wait for a **specific child** to terminate.
3. **Q:** Difference between wait() and waitpid()?  
**A:**
  - wait() → waits for **any child**.
  - waitpid() → waits for a **specific child** by PID.
4. **Q:** Why do we use exit(0) in child processes?  
**A:** To terminate the child cleanly and indicate successful completion.
5. **Q:** What is WIFEXITED(status) and WEXITSTATUS(status)?  
**A:**
  - WIFEXITED(status) → checks if child exited normally.
  - WEXITSTATUS(status) → returns the exit code of the child.
6. **Q:** Can parent wait for multiple children with waitpid()?  
**A:** Yes, by calling waitpid() separately for each child's PID.
7. **Q:** What happens if we use wait(NULL) instead of waitpid()?  
**A:** Parent waits for **any child** to terminate, not a specific one; output order may vary.

8. **Q:** Why is sleep(2) used in children?

**A:** To simulate processing time and make child execution visible; helps observe output order.

9. **Q:** What is the type returned by fork()?

**A:** pid\_t (integer type representing process IDs).

10. **Q:** Can fork() fail? How do you handle it?

**A:** Yes, if the system is out of resources. Handle using:

**Q26. Write a program to execute an executable program. (a) use some executable program (b) pass some input to an executable program.(for example execute an executable of `./a.out` name).**

**Code Explanation**

```
#include <stdio.h>
#include <unistd.h>
```

- stdio.h → for standard input/output functions like printf() (though not used here).
- unistd.h → for **exec()** and other POSIX system calls.

---

```
int main(int argc, char *argv[])
```

- argc → counts the **number of command-line arguments**.
- argv → array of strings containing **command-line arguments**.
  - argv[0] → name of the program (26 if run as ./26).
  - argv[1], argv[2] → first and second arguments passed to program.

---

```
exec("./function", "function", argv[1], argv[2], NULL);
```

- **exec()** is a system call that **replaces the current process** with a new program.
- Arguments:
  1. "./function" → path to executable to run.
  2. "function" → argv[0] for the new program (can be the program name).
  3. argv[1], argv[2] → arguments passed from current program.
  4. NULL → marks the end of argument list (mandatory).

**Important points:**

- After `exec()`, the current program **does not continue** if successful.
  - If `exec()` fails (e.g., file not found), it **returns -1**.
- 

```
return 0;
```

- This runs **only if exec() fails**, since a successful `exec()` does **not return**.
- 

## Example Execution

```
gcc 26.c -o 26
```

```
./26 hello world
```

- The program will attempt to run `./function hello world`.
  - If `./function` exists, it **replaces the current process**, so nothing else in `26.c` runs.
  - If it fails, program returns 0 (or better to handle failure with `perror()`).
- 

## Key Viva Questions

1. **Q:** What is `exec()`?  
**A:** It replaces the current process image with a new program, passing arguments as a list.
2. **Q:** How is `exec()` different from `fork()`?  
**A:**
  - `fork()` → creates a **new child process** that runs concurrently.
  - `exec()` → **replaces the current process**, does not create a child.
3. **Q:** What happens after a successful `exec()`?  
**A:** The current process is replaced; **no code after exec() executes**.
4. **Q:** Why is `NULL` required at the end of `exec()`?  
**A:** To indicate the **end of arguments**; otherwise, `exec()` doesn't know where the list ends.
5. **Q:** What are `argv[1]` and `argv[2]`?  
**A:** Command-line arguments passed to the program; in this case, forwarded to `./function`.
6. **Q:** What happens if `./function` does not exist?  
**A:** `exec()` fails and returns -1. Program continues (or should print an error with `perror()`).
7. **Q:** How is `exec()` different from `execv()`?  
**A:**
  - `exec()` → arguments passed as **list**.
  - `execv()` → arguments passed as **array of strings** (`char* argv[]`).

8. **Q:** Does the program create a new process?

**A:** No, it **replaces the current process**. If you need a new process, you first call fork() then execl() in the child.

9. **Q:** Why might you include perror("execl failed")?

**A:** To **debug** and see why execl() failed (e.g., file not found or permission denied).

## Q27. Write a program to execute ls -Rl by the following system calls (a) execl (b) execvp (c) execle (d) execv (e) execvp

**Code Explanation:**

```
#include <stdio.h>
#include <unistd.h>
```

- unistd.h → contains all the **exec family system calls** (execl, execv, etc.)
- stdio.h → standard I/O (not used here but commonly included)

---

```
int main(int argc, char *argv[])
{
```

- Standard main function with command-line arguments, though **not used** in this program.

---

### 1. execl()

```
execl("/bin/ls", "ls", "-Rl", NULL);
```

- **Replaces the current process** with /bin/ls.
- Arguments passed as a **list**:
  1. /bin/ls → path to executable
  2. ls → argv[0] (program name)
  3. -Rl → argument to ls
  4. NULL → end of list
- If successful, **code after this line is never executed**.

## 2. execp()

```
execp("ls", "ls", "-Rl", NULL);
```

- Similar to execl(), but **searches PATH** environment variable to locate ls.
  - You don't need to give the full path.
- 

## 3. execle()

```
char *envp[] = {"PATH=/bin", NULL};  
execle("/bin/ls", "ls", "-Rl", NULL, envp);
```

- execle() → similar to execl(), but allows you to **specify a custom environment** (envp).
  - envp[] defines **environment variables** for the new process.
- 

## 4. execv()

```
char *args[] = {"ls", "-Rl", NULL};  
execv("/bin/ls", args);
```

- execv() → arguments passed as **array** instead of a list.
  - First element is argv[0] (program name), last element must be NULL.
- 

## 5. execvp()

```
execvp("ls", args);
```

- execvp() → like execv(), but **searches PATH** to locate the program.
  - Handy when you don't know the full path.
- 

## Important Notes

1. **Only the first exec that succeeds runs.**
    - After a successful exec\*(), the current process is replaced; no code after it executes.
  2. **If exec fails**, control moves to the next line.
    - Usually you should add perror("exec failed") to catch errors.
  3. return 0; is executed only if **all exec calls fail**.
- 

## Viva Questions

1. **Q:** What is the exec family of functions?

**A:** System calls that **replace the current process** with a new program.

2. **Q:** What's the difference between execl and execv?

**A:**

- o execl → arguments as a **list**
- o execv → arguments as an **array**

3. **Q:** Difference between execl and execlp?

**A:** execlp searches the **PATH** environment to locate the executable; execl needs the **full path**.

4. **Q:** Difference between execle and execl?

**A:** execle allows you to **specify a custom environment** for the new process.

5. **Q:** Difference between execv and execvp?

**A:** execvp searches PATH to find the executable; execv requires **full path**.

6. **Q:** What happens if an exec call succeeds?

**A:** The **current process is replaced**, code after the exec **does not execute**.

7. **Q:** What happens if an exec call fails?

**A:** It returns -1, and you can use perror() to print the error.

8. **Q:** Can you run multiple exec calls sequentially like this?

**A:** Only if **previous exec calls fail**, because a successful exec replaces the process.

9. **Q:** Why is NULL required at the end of arguments or env array?

**A:** It **terminates the argument or environment list**, required by the kernel.

## Q28. Write a program to get maximum and minimum real time priority.

Code Explanation: sched\_get\_priority\_min/max

```
#include <stdio.h>
```

```
#include <sched.h>
```

- sched.h → contains functions and macros for **real-time scheduling** in Linux (SCHED\_FIFO, SCHED\_RR, etc.)

```
int main()
{
    int mpf = sched_get_priority_min(SCHED_FIFO);
    int mxpf = sched_get_priority_max(SCHED_FIFO);

    int mpr = sched_get_priority_min(SCHED_RR);
```

```
int mxpr = sched_get_priority_max(SCHED_RR);
```

- `sched_get_priority_min(policy)` → returns **minimum priority** allowed for the given scheduling policy.
- `sched_get_priority_max(policy)` → returns **maximum priority** allowed.
- `SCHED_FIFO` → **First-In-First-Out** real-time scheduling policy.
- `SCHED_RR` → **Round-Robin** real-time scheduling policy.
- Example output (depends on system):
  - 1 99
  - 1 99
    - Real-time priorities in Linux usually range **1–99**.

```
printf("%d %d\n", mpf, mxpf);
printf("%d %d\n", mpr, mxpr);

return 0;
}
```

- Prints **minimum and maximum priorities** for both `SCHED_FIFO` and `SCHED_RR`.

## Important Notes

1. These functions **do not set the priority**; they only query the system limits.
2. Only **root users** can set real-time scheduling policies.

## Viva Questions

1. **Q:** What is `sched_get_priority_min()`?

**A:** Returns the **minimum priority** allowed for a specific scheduling policy.

2. **Q:** What is `sched_get_priority_max()`?

**A:** Returns the **maximum priority** allowed for a specific scheduling policy.

3. **Q:** What are `SCHED_FIFO` and `SCHED_RR`?

**A:**

- `SCHED_FIFO` → First-In-First-Out real-time scheduling
- `SCHED_RR` → Round-Robin real-time scheduling

4. **Q:** Can normal users change the priority of a process?  
**A:** No, **only root** can assign real-time priorities.
5. **Q:** What is the typical priority range in Linux for real-time policies?  
**A:** 1 (minimum) to 99 (maximum)
6. **Q:** What happens if you pass an invalid policy to sched\_get\_priority\_min/max?  
**A:** It returns -1 and sets errno.
7. **Q:** Difference between SCHED\_OTHER, SCHED\_FIFO, and SCHED\_RR?  
**A:**
  - o SCHED\_OTHER → normal time-sharing scheduling
  - o SCHED\_FIFO → first-in-first-out real-time
  - o SCHED\_RR → round-robin real-time

## Q29. Write a program to get scheduling policy and modify the scheduling policy (SCHED\_FIFO, SCHED\_RR).

### Code Explanation

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sched.h>
```

- sched.h → provides **scheduling functions and structures** (sched\_getscheduler, sched\_setscheduler, sched\_param, etc.)
- unistd.h → system calls
- stdio.h/stdlib.h → standard I/O and utilities

### Function to Print Policy

```
void print_policy(int policy){
    if (policy == SCHED_OTHER)
        printf("Policy: SCHED_OTHER (normal)\n");
    else if (policy == SCHED_FIFO)
        printf("Policy: SCHED_FIFO\n");
    else if (policy == SCHED_RR)
        printf("Policy: SCHED_RR\n");
```

```

else
    printf("Policy: UNKNOWN (%d)\n", policy);
}

```

- Converts a **numeric policy** into a **readable string**.
  - Useful for displaying results.
- 

## Main Function

```

int main() {
    int pid = 0;
    int old_policy, new_policy;
    struct sched_param param;
}

```

- pid = 0 → refers to **current process** in sched\_getscheduler and sched\_setscheduler.
  - struct sched\_param param → holds **priority** for real-time scheduling.
- 

## Get Current Scheduling Policy

```

old_policy = sched_getscheduler(pid);
printf("Old ");
print_policy(old_policy);
}

```

- sched\_getscheduler(pid) → returns **current scheduling policy** for the process.
  - Prints old/current policy using print\_policy().
- 

## Set New Scheduling Policy

```

param.sched_priority = 10;
sched_setscheduler(pid, SCHED_FIFO, &param);
}

```

- Sets **new scheduling policy** to SCHED\_FIFO (real-time FIFO).
  - sched\_priority = 10 → priority for the real-time process (must be between min/max for SCHED\_FIFO).
  - **Root privileges required** to set real-time scheduling.
- 

## Check New Policy

```

new_policy = sched_getscheduler(pid);
}

```

```
printf("New ");  
print_policy(new_policy);
```

- Verifies that the scheduling policy has changed successfully.
- 

## Return Statement

```
return 0;
```

- Program ends successfully.
- 

## Important Notes

1. Only **root** can set real-time scheduling policies (SCHED\_FIFO, SCHED\_RR).
  2. SCHED\_OTHER → default Linux time-sharing policy for normal processes.
  3. sched\_getscheduler() returns -1 on failure (check errno).
  4. sched\_setscheduler() can fail if priority is **out of range** or not root.
- 

## Viva Questions

1. **Q:** What is sched\_getscheduler()?

**A:** Returns the **current scheduling policy** of a process.

2. **Q:** What is sched\_setscheduler()?

**A:** Sets a **new scheduling policy** and priority for a process.

3. **Q:** What is struct sched\_param?

**A:** Structure containing scheduling parameters like sched\_priority.

4. **Q:** What are the scheduling policies in Linux?

**A:**

- SCHED\_OTHER → normal time-sharing
- SCHED\_FIFO → first-in-first-out real-time
- SCHED\_RR → round-robin real-time

5. **Q:** What happens if a normal user tries to set SCHED\_FIFO?

**A:** sched\_setscheduler() fails with **EPERM** (permission denied).

6. **Q:** What is the valid priority range for SCHED\_FIFO?

**A:** sched\_get\_priority\_min(SCHED\_FIFO) to sched\_get\_priority\_max(SCHED\_FIFO) (usually 1–99).

7. **Q:** What is the difference between SCHED\_FIFO and SCHED\_RR?

**A:**

- SCHED\_FIFO → executes in **first-come, first-served** order
- SCHED\_RR → executes in **round-robin slices**

8. **Q:** What is the significance of pid = 0?

**A:** Refers to the **calling process itself**.

### Q30. Write a program to run a script at a specific time using a Daemon process.

**Code Explanation: Creating a Daemon Process**

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <time.h>
#include <signal.h>
#include <syslog.h>
```

- Includes necessary headers for process management, signals, logging, and file system operations.

```
#define LOGGING "Start Logging my task = %d\n"
```

- A macro for the log message format.

**Step 1: First fork**

```
pid = fork();
if (pid > 0)
    exit(EXIT_SUCCESS);
```

- Creates a **child process**.
- Parent exits (exit(EXIT\_SUCCESS)), allowing the child to **run independently**.

**Step 2: Create new session**

```
if (setsid() < 0)
    exit(EXIT_FAILURE);
```

- setsid() creates a **new session**:
    - The child becomes **session leader**
    - Detached from controlling terminal
    - Becomes **leader of a new process group**
- 

### Step 3: Ignore signals

```
signal(SIGCHLD, SIG_IGN);  
signal(SIGHUP, SIG_IGN);
```

- SIGCHLD → ignore termination of child processes.
  - SIGHUP → ignore hangup signal (terminal closes).
- 

### Step 4: Second fork

```
pid = fork();  
  
if (pid > 0)  
{  
    printf("Daemon PID: %d\n", pid);  
    exit(EXIT_SUCCESS);  
}
```

- Second fork ensures **daemon is not session leader** → cannot accidentally acquire a controlling terminal.
- 

### Step 5: Set file permissions and working directory

```
umask(077);  
chdir("/");
```

- umask(077) → new files are **accessible only by owner**.
  - chdir("/") → sets working directory to root, avoiding locking a mounted filesystem.
- 

### Step 6: Close open file descriptors

```
for (x_fd = sysconf(_SC_OPEN_MAX); x_fd >= 0; x_fd--)  
  
close(x_fd);  
  
• Closes all inherited file descriptors (stdin, stdout, stderr, etc.)
```

---

## Step 7: Logging loop

```
int count = 0;  
  
openlog("Logs", LOG_PID, LOG_USER);  
  
while (1)  
{  
    sleep(2);  
    syslog(LOG_INFO, LOGGING, count++);  
}  
  
closelog();
```

- `openlog()` → initialize connection to syslog with identifier "Logs".
  - Infinite loop logs a **counter every 2 seconds** using `syslog(LOG_INFO, ...)`.
  - `closelog()` → closes connection to syslog (though never reached in infinite loop).
- 

## Key Concepts

1. **Daemon process** → background process detached from terminal.
  2. **Double fork technique** → ensures daemon cannot acquire controlling terminal.
  3. **syslog** → used to log messages to system log.
  4. **umask, chdir, close file descriptors** → best practices for daemons.
- 

## Viva Questions

1. **Q:** What is a daemon process?  
**A:** A background process that runs independently of the terminal.
2. **Q:** Why do we use double fork to create a daemon?  
**A:** First fork → child runs in background. Second fork → ensures **daemon cannot acquire controlling terminal**.
3. **Q:** What does setsid() do?  
**A:** Creates a **new session**, makes the process **session leader** and detaches it from terminal.
4. **Q:** Why do we ignore SIGHUP and SIGCHLD?  
**A:**
  - SIGHUP → prevents daemon termination when terminal closes
  - SIGCHLD → prevents zombie processes when child terminates

5. **Q:** Why set umask(077)?

**A:** Ensures newly created files are **accessible only by owner**, for security.

6. **Q:** Why change working directory to /?

**A:** Avoids locking the current directory, especially if it's on a **mounted filesystem**.

7. **Q:** Why close all file descriptors?

**A:** Daemon should **not use stdin, stdout, stderr**; avoids interfering with terminal or files.

8. **Q:** What is syslog used for?

**A:** To log messages to the **system log** instead of console.

9. **Q:** How often does this daemon log messages?

**A:** Every 2 seconds, increments a counter.

10. **Q:** What privileges are required to create a daemon?

**A:** Normal user can create a daemon; **root privileges** may be needed for certain syslog operations or file permissions.